

RF-100 — Verified Execution Governance Standard

Remnant Fieldworks Inc. — Standards Track

Version 1.0 — Unified DRAFT for Release (second-pass corrected)

Document status: DRAFT v1.0 (unified), second-pass corrected. Supersedes RF-100 v0.9 and the interim releases v1.0-draft-g18-corrected and v1.0-draft Technical Gap Integration, and the July 16, 2026 unified draft submitted for second-pass review.

Incorporates: full disposition (Accept) of Gap Register G-01 through G-17 and the G-18 vocabulary/posture correction from the Structured Peer Review by James A. Bex, BxTech Systems Solutions LLC (July 15, 2026); and the six findings (F-01 through F-06) of the Second-Pass Review Memorandum by James A. Bex (July 16, 2026). See Annex B for the change log and Annex E for the second-pass disposition record.

Publisher: Remnant Fieldworks Inc. · Derek Hone, Founder

Date: July 2026

Prior version of record: RF-100 v0.9, Zenodo DOI 10.5281/zenodo.21366342 · github.com/derekhone/rf-100

Contents

- 1. Status, Conformance Language, and Notation
- 1. Purpose and Scope
- 1. Terms, Definitions, and Roles
- 1. Policy and Execution Boundary Profile (EBP)
- 1. Evidence
- 1. Verification and Decision Inputs
- 1. Authority
- 1. Decisions: Rendering, Validity, and Failure Posture
- 1. Verification Decision Records (VDRs)
- 1. Execution Integrity
- 1. Auditability and Reconstruction
- 1. Enforcement Architecture
- 1. Gate Governance
- 1. Trusted Time
- AI. AI Actor Profile (RF-100:AI)
- Annex A — Composed Decision Mechanism (Informative)
- Annex B — Change Log v0.9 → v1.0 (Informative)
- Annex C — Adversary Model (Normative)
- Annex D — Conformance Claims and RF-101 (Normative)

0. Status, Conformance Language, and Notation

The key words **MUST**, **MUST NOT**, **SHALL**, **SHALL NOT**, **SHOULD**, **SHOULD NOT**, and **MAY** are to be interpreted as in RFC 2119/8174. Requirements are individually numbered. Each requirement carries a verification method tag: **[T]** Test, **[D]** Demonstration, **[I]** Inspection, **[A]** Analysis.

This standard is implementation-neutral. Conformance with RF-100 does not require the use of any Remnant Fieldworks product, service, trademark, hosted endpoint, or branded implementation. The normative vocabulary of this standard is generic: **Verification Gate**, **Verification Decision Record (VDR)**, **Governed Action**, **Execution Boundary Profile (EBP)**, and **Enforcement Point**. ExecutionProof™ (commercial platform), ProofRecord™ (a branded implementation of the Verification Decision Record), Proof Before Power™, and Verification Before Execution™ are Remnant Fieldworks brands and appear nowhere in normative language.

0.1 Licensing and IPR Notice

Specification Text License. The RF-100 specification text is intended to be made available under CC BY 4.0 upon v1.0 release. This license applies only to the text of the RF-100 specification and does not grant any license to Remnant Fieldworks trademarks, service marks, certification marks, software, hosted services, APIs, trade secrets, or patents.

Patent/IPR Position. Remnant Fieldworks supports open implementation and independent assessment of RF-100. A **draft IPR commitment** describing the intended patent posture — a royalty-free, scoped, reciprocal (defensive) non-assertion covenant for conforming implementations, with rights reserved for trademarks, products, hosted services, reference code, and the certification program — is published alongside this draft as `IPR-COMMITMENT.md` in the RF-100 repository. That document (v0.2-draft) has been revised to incorporate a pre-counsel legal review — adding successors/assigns and runs-with-the-patents language, a version-lock irrevocability rule, third-party-beneficiary intent, defined terms, worldwide scope, an Ohio governing-law clause, and a no-implied-rights clause — but remains a **draft pending formal legal/patent counsel review** (see its Annex IPR-A for open action items). The binding version will be finalized and published before any formal certification or conformance program (RF-101 / RF-102) is launched.

1. Purpose and Scope

RF-100 defines requirements for governing consequential actions — Governed Actions — proposed by human or AI Actors, such that no Governed Action takes effect without independent verification of authority, evidence, policy compliance, and system state, and such that every decision is durably and verifiably recorded.

RF-100 applies to any system in which Actors (human, automated, or AI agent) propose actions with real-world consequence: financial transactions, infrastructure changes, data operations, physical actuation, communications, and administrative operations — including administration of the governance system itself (Section 12).

The central architectural premise, stated plainly: the proposing Actor is precisely the component that cannot be trusted to self-govern. Verification and enforcement therefore reside outside the Actor's control (Sections 5, 11, Annex C).

2. Terms, Definitions, and Roles

Term	Definition
Actor	An entity (human, automated system, or AI agent) that proposes a Governed Action.
Governed Action	An action designated by the EBP as requiring verification before effect.
Verification Gate (“gate”)	The component that evaluates a proposed Governed Action against policy, evidence, authority, and state, and renders a Decision.
Decision	The rendered outcome of verification: ALLOW, HOLD, or DENY.
Verification Decision Record (VDR)	The signed, durable record of a Decision and its complete basis (Section 8).
Execution Boundary Profile (EBP)	The declared, versioned profile defining the Governed Action set, thresholds, evidence requirements, authority model, enforcement class, availability class, validity windows, and consistency model for a deployment (Section 3).
Evidence	Integrity-protected artifacts on which a Decision is based (Section 4).
Intent Record	Actor-asserted statement of purpose accompanying a proposal. Untrusted context, never Evidence (AI.4a).
Enforcement Point	A component outside the proposing Actor’s administrative control that permits or blocks execution based on an authenticated Decision.
Effector	The component that produces the real-world effect of a Governed Action.
Consequence Threshold	The EBP-declared boundary above which elevated requirements (E2+, AI.7, AI.10, AI.11) apply.
Delegation Chain	The ordered record of authority from the originating authority through any intermediary Actors to the proposing Actor, including, in chained or multi-agent workflows, each consequential hop (AI.8a, 8.1).

3. Policy and Execution Boundary Profile (EBP)

RF-100.3.1 — Each deployment **MUST** declare a versioned EBP enumerating: the Governed Action set; consequence tiers and the consequence threshold; evidence requirements and freshness limits per action class; the authority model; the Enforcement Class per Governed Action (Section 11); the Gate Availability Class (7.8); decision validity windows (7.6); and, where applicable, the velocity-limit consistency model (AI.6a). **[I]**

RF-100.3.2 — Policy **MUST** be versioned and content-addressable. The policy applied to any Decision is identified by cryptographic content hash, not by label alone. **[I]**

RF-100.3.3 — Implementations **MUST** support evaluation of policy in simulation (shadow) mode without producing execution effects. **[D]**

RF-100.3.4 — Each Decision **MUST** be evaluated in its entirety against exactly one immutable policy snapshot, identified by cryptographic content hash. The snapshot hash **MUST** be recorded in the VDR. Policy updates **MUST NOT** alter the snapshot applied to a Decision already in evaluation. **[T]**

RF-100.3.5 — Implementations **SHOULD** employ drain-and-swap policy rollout (in-flight Decisions complete against the prior snapshot before a new snapshot activates) and **SHOULD** evaluate new policy snapshots in shadow against live traffic before enforcement. **[D]**

4. Evidence

RF-100.4.1 — Every Governed Action class in the EBP **MUST** specify the Evidence required for an ALLOW. **[I]**

RF-100.4.2 — Evidence **MUST** carry a declared freshness limit; Evidence older than its limit **MUST NOT** support an ALLOW. **[T]**

RF-100.4.3 — Absence of required Evidence **MUST** result in HOLD or DENY, never ALLOW. **[T]**

RF-100.4.4 — Evidence artifacts for Governed Actions **MUST** be integrity-protected from source to gate, and the evidence source **MUST** be authenticated (signed at origin, or collected over an authenticated channel with source identity recorded). Evidence references in the VDR **MUST** include a cryptographic digest of each artifact as evaluated. **[T]**

5. Verification and Decision Inputs

RF-100.5.1 — Verification **MUST** be performed by a Verification Gate that is administratively and operationally independent of the proposing Actor. The Actor **MUST NOT** be able to modify gate policy, gate configuration, Evidence in transit, or the Decision. **[A]**

RF-100.5.2 — Every Decision **MUST** be based on: (a) verified authority of the Actor (Section 6); (b) required Evidence (Section 4); (c) the pinned policy snapshot (3.4); and (d) evaluated system state and risk signals as declared in the EBP. **[T]**

RF-100.5.3 — Decision outcomes are ALLOW, HOLD, or DENY. No other outcome **SHALL** authorize execution. **[I]**

RF-100.5.4 — Only artifacts satisfying Section 4 constitute Evidence. Actor-asserted context (including Intent Records) MUST NOT be treated as Evidence. **[I]**

6. Authority

RF-100.6.1 — Authority to propose or approve each Governed Action class MUST be explicitly bound to identities in the EBP authority model. Implicit or ambient authority MUST NOT support an ALLOW. **[I]**

RF-100.6.2 — Where M-of-N approval is required, approvals MUST be independently authenticated and recorded in the VDR. Approver independence for AI Actors is governed by AI.11. **[T]**

RF-100.6.3 — Authority revocation MUST take effect for all Decisions not yet rendered; interaction with in-flight evaluation is governed by the pinning rule (3.4) and validity window (7.6), such that no ALLOW outlives both its snapshot's validity and its window. **[T]**

7. Decisions: Rendering, Validity, and Failure Posture

RF-100.7.1 — A Decision MUST be rendered before any Governed Action takes effect. **[T]**

RF-100.7.2 — The system MUST fail closed: under any condition enumerated in the assessment fault matrix (Annex D.2) — gate outage, network partition, evidence-service failure, authority-service failure, clock fault — absence of a valid Decision MUST result in no execution. **[T]**

RF-100.7.3 — HOLD MUST route the proposal to the escalation path declared in the EBP. **[D]**

RF-100.7.4 — An unverifiable, malformed, unauthenticated, expired, or already-consumed Decision MUST be treated as absent (see 7.10). **[T]**

RF-100.7.5 — DENY MUST be recorded with the same completeness as ALLOW (Section 8). **[I]**

RF-100.7.6 — Each EBP MUST define a decision validity window per Governed Action class. A Decision whose validity window has elapsed MUST be treated as absent; execution against an expired ALLOW MUST NOT proceed. **[T]**

RF-100.7.7 — Resolution of a HOLD MUST trigger re-verification of authority, evidence freshness, system state, and risk before an ALLOW is rendered. Approval of a held action authorizes re-evaluation, not execution. **[T]**

RF-100.7.8 — Each EBP MUST state a Gate Availability Class commensurate with the consequence tier of the Governed Actions it gates, and MUST define the disposition of in-flight HOLDS during gate unavailability. **[I]**

RF-100.7.9 — Where emergency execution during gate unavailability is operationally required, the EBP MAY define pre-authorized break-glass instruments: single-use, action-bound, time-bound tokens issued in advance as Governed Actions. Break-glass use MUST generate a VDR upon gate restoration and MUST trigger mandatory review. No other mechanism SHALL permit execution during gate unavailability. **[D]**

RF-100.7.10 — Prior to producing any execution effect, the executing component or its Enforcement Point MUST authenticate the Decision: verifying the gate signature (8.4), the action parameter digest

(8.1), the single-use status (9.3), and the validity window (7.6). A Decision that fails any check **MUST** be treated as absent. **[T]**

8. Verification Decision Records (VDRs)

RF-100.8.1 — Every Decision **MUST** produce a VDR containing at minimum: the canonical action identifier and a cryptographic digest of the complete canonicalized parameter set of the proposed action; Actor identity (per AI.1 for AI Actors); the Decision and its basis; the policy snapshot hash (3.4); Evidence references with content digests (4.4); authority and approval records (6.2), **including the complete delegation chain from originating authority to proposing Actor**; gate identity including gate software version and configuration hash; timestamps (Section 13); and the validity window (7.6). **[I]**

RF-100.8.2 — VDRs **MUST** be durable and retained per the EBP retention schedule. **[I]**

RF-100.8.3 — VDRs **MUST** be independently verifiable offline by parties other than the record custodian. **[T]**

RF-100.8.4 — VDRs **MUST** be signed with an asymmetric digital signature scheme drawn from a documented, current national or international standard (current baseline: Ed25519; post-quantum target: ML-DSA / FIPS 204). Symmetric message authentication codes **MUST NOT** be the sole integrity mechanism for conformance claims. **[T]**

RF-100.8.5 — Implementations **MUST** document a cryptographic agility plan, including a migration path to standardized post-quantum signature schemes, covering all records within their retention windows. **[I]**

RF-100.8.6 — VDRs **SHOULD** be committed to an append-only, externally anchored transparency structure (e.g., Merkle inclusion proofs anchored periodically) such that deletion or modification of any record is detectable by parties other than the record custodian. **[D]**

RF-100.8.7 — VDR stores **MUST** enforce documented access control and confidentiality protections commensurate with record sensitivity. EBPs **MUST** apply field-level minimization (record what is needed to reconstruct the decision, not the maximum available). Where legally mandated erasure applies, implementations **SHOULD** support cryptographic erasure (crypto-shredding) of protected fields without invalidating chain integrity of remaining records. **[I]**

RF-100.8.8 — Gate signing keys **MUST** be protected in an HSM or equivalent isolation. Signing-key generation and rotation are Governed Actions under Section 12. **[I]**

9. Execution Integrity

RF-100.9.1 — Execution **MUST** correspond to the action as verified. **[A]**

RF-100.9.2 — At Enforcement Class E2 and above, divergence between the verified action and the executed action **MUST** be prevented and detected per 9.2a. At Class E1, the EBP **MUST** declare a post-hoc reconciliation procedure comparing effector records against VDR parameter digests on a stated cadence; divergences found **MUST** be recorded and alerted. **[T]**

RF-100.9.2a — At Enforcement Class E2 or higher, the Enforcement Point MUST verify that the parameter digest of the action being executed matches the digest recorded in the authorizing VDR; on mismatch, execution MUST NOT proceed and the event MUST be recorded and alerted. **[T]**

RF-100.9.3 — A Decision MUST be bounded in reuse; unless the EBP explicitly declares a bounded multi-use window for a specific action class, each ALLOW is single-use. **[T]**

10. Auditability and Reconstruction

RF-100.10.1 — Every Governed Action outcome (executed, denied, held, expired, break-glass) MUST be traceable to its VDR. **[I]**

RF-100.10.2 — Audit access MUST NOT permit modification of VDRs. **[T]**

RF-100.10.3 — It MUST be possible to reconstruct, from the VDR and referenced artifacts alone, the complete basis of any Decision: who proposed, under what authority, on what evidence, against which policy snapshot, with what outcome, enforced how. **[D]**

11. Enforcement Architecture

Class	Name	Definition	Permitted use
E1	Cooperative	The Actor invokes the gate via SDK/API and honors the Decision voluntarily.	Low-consequence actions only. MUST NOT be claimed for AI-profile actions at/above the consequence threshold (AI.10).
E2	Interposed	The gate or an Enforcement Point it controls sits in the execution path (API gateway, tool-call proxy, MCP gateway, service mesh); no invocation reaches the Effector without a rendered Decision.	Minimum class for AI-profile actions at/above the consequence threshold.
E3	Infrastructure-enforced	The ability to execute is itself gated: execution credentials, capability tokens, or network reachability are issued only upon ALLOW; bypass requires compromising infrastructure.	Highest-consequence actions.

RF-100.11.1 — Each EBP MUST declare the Enforcement Class at which the Verification Gate is bound to the execution pathway, per Governed Action class, and every conformance claim MUST state the Enforcement Class per Governed Action. **[I]**

RF-100.11.2 — At Class E2 and above, no Governed Action invocation SHALL be able to reach its Effector without a Decision having been rendered; the Enforcement Point MUST be outside the administrative control of the proposing Actor. **[T]**

12. Gate Governance

RF-100.12.1 — Gate administration actions — including gate software deployment, configuration changes, signing-key operations (8.8), EBP modification, and approver-set changes — are Governed Actions and MUST be verified prior to taking effect, by a Verification Gate instance distinct from the one being administered or by a documented M-of-N administrative control. Each such action MUST generate a VDR. **[D]**

13. Trusted Time

RF-100.13.1 — The Verification Gate **MUST** synchronize to an authenticated time source, with a documented maximum drift no greater than a stated fraction of the tightest Evidence Freshness limit or decision validity window it enforces. The drift bound **MUST** be stated in the EBP. VDR timestamps **SHOULD** be independently verifiable (e.g., RFC 3161 timestamping or anchoring per 8.6). **[T]**

RF-100:AI — AI Actor Profile

This profile applies whenever the proposing Actor is, or materially incorporates, an AI system. Its premise (Annex C): the agent is precisely the component that cannot be trusted to self-govern.

RF-100:AI.1 — Agent identity **MUST** include model/version, base-model lineage, and deployment context, recorded in every VDR the agent's proposals generate. **[I]**

RF-100:AI.2 — Agent capability **MUST** be scoped: an agent **MAY** propose only Governed Actions within its EBP-declared capability set. **[T]**

RF-100:AI.3 — Agent-proposed Governed Actions **MUST** pass through the same gate, evidence, and authority requirements as human-proposed actions; no reduced path exists for automation. **[I]**

RF-100:AI.4 — Agents **SHOULD** submit Intent Records with proposals to support consistency evaluation and audit. **[I]**

RF-100:AI.4a — Intent Records are actor-asserted context. They **MUST NOT** be treated as Evidence under Section 5.4 and **MUST NOT** be a basis for elevating a decision toward ALLOW. Consistency evaluation outputs **MAY** only contribute risk signals toward HOLD or DENY. Any model-based consistency evaluator **MUST** be isolated from the proposing Actor's input context. Deterministic checks (parameter schema, scope, threshold comparisons) gate the decision. **[A]**

RF-100:AI.5 — The EBP **MUST** define the consequence threshold above which elevated AI-profile requirements (AI.7, AI.10, AI.11) apply. **[I]**

RF-100:AI.6 — Cumulative and velocity limits **MUST** be defined for agent-proposed Governed Action classes and evaluated at the gate. **[T]**

RF-100:AI.6a — Cumulative and velocity evaluations **MUST** be computed against strongly consistent shared state across all gate instances, or against statically partitioned per-instance quotas whose sum does not exceed the global limit. The EBP **MUST** document the consistency model employed. **[T]**

RF-100:AI.7 — Prompt- or input-injection indicators **MUST** be treated as risk signals for Governed Actions at or above the consequence threshold. Detection **MUST** operate outside the proposing Actor's observable or influenceable context. A flagged session **MUST** remain flagged for its duration, and above-threshold actions proposed within it **MUST NOT** receive ALLOW without elevated review. **[T]**

RF-100:AI.8 — Agent sessions **MUST** be attributable: every proposal is bound to a session identity traceable across the VDR chain. **[I]**

RF-100:AI.8a — In chained or multi-agent workflows, each consequential hop that meets the Governed Action definition is independently a Governed Action, subject to full verification, with the com-

plete delegation chain from the originating Actor preserved in each VDR. An ALLOW rendered for one hop MUST NOT authorize any subsequent hop. **[T]**

RF-100:AI.9 — Agents MUST NOT hold, issue, or approve their own authority; authority binding is external to the agent (Section 6). **[A]**

RF-100:AI.10 — Agent-proposed Governed Actions at or above the consequence threshold MUST be enforced at Class E2 or higher. Class E1 MUST NOT be claimed as conformant for such actions. **[I]**

RF-100:AI.11 — For Governed Actions at or above the consequence threshold, an M-of-N approver set composed of AI Actors MUST include at least one human approver or at least two approvers of distinct base-model lineage. Approvers sharing base-model lineage SHALL count as one approver for independence purposes. **[I]**

RF-100:AI.12 — A change to an agent’s model, base-model lineage, weights (including fine-tuning), or system prompt constitutes a new agent identity version under AI.1 and MUST trigger re-evaluation of the agent’s EBP scope, capability set, consequence-threshold assignments, and velocity limits before the updated agent proposes further Governed Actions. **[I]**

Annex A — Composed Decision Mechanism (Informative)

A conforming Decision at Class E2 or higher is:

Property	Mechanism	Requirement
Signed	Asymmetric gate signature	8.4
Parameter-bound	Canonical action digest in VDR; recheck at effect	8.1, 9.2a
Time-bound	Decision validity window (TTL)	7.6
Single-use	Bounded reuse	9.3
Enforced beyond actor control	Enforcement Class E2/E3	11.1, 11.2, AI.10
Authenticated at effect	Executor verification of signature, digest, use, TTL	7.10
Policy-pinned	Immutable snapshot hash	3.4
Evidence-bound	Source-authenticated, digest-referenced artifacts	4.4

Annex B — Change Log v0.9 → v1.0 (Informative)

Gap	Sev.	Resolution in v1.0
G-01	Critical	Section 11 (E1/E2/E3); 11.1, 11.2; AI.10
G-02	Critical	8.1 amended (action ID + parameter digest); 9.2a
G-03	High	7.6 (validity window); 7.7 (post-HOLD re-verification)
G-04	High	7.8 (availability class, HOLD disposition); 7.9 (break-glass)
G-05	High	3.4 (decision pinning); 3.5 (rollout guidance)
G-06	High	AI.1 amended (base-model lineage); AI.11
G-07	High	8.3 SHOULD→MUST; 8.4, 8.5, 8.6, 8.8
G-08	High	Section 12; 12.1; 8.1 amended (gate version + config hash)
G-09	High	7.10 (executor authentication of decisions)
G-10	Medium	Normative Annex C (adversary model)
G-11	Medium	AI.4a (Intent Records untrusted; evaluator isolation); 5.4
G-12	Medium	AI.7 amended (MUST above threshold; out-of-band; flag persistence)
G-13	Medium	4.4 amended to MUST (provenance, source auth, digests)
G-14	Medium	Section 13; 13.1 (trusted time)
G-15	Medium	8.7 (confidentiality, minimization, crypto-shredding)

Gap	Sev.	Resolution in v1.0
G-16	Medium	AI.6a (consistency model for velocity limits)
G-17	Medium	Verification method tags on all requirements; fault matrix (Annex D.2); skeletal RF-101 companion
G-18	High	Partially resolved. Vocabulary and neutrality complete (generic VDR terminology; de-trademarked normative text; implementation-neutral conformance). CC BY 4.0 (spec text only) declared for v1.0 and effective at release. Patent/IPR commitment now published in draft (IPR-COMMITMENT.md , Section 0.1): royalty-free, scoped, reciprocal non-assertion for conforming implementations, rights reserved for brands/products/services/cert. The draft is pending counsel review; the binding version will be finalized before any certification program launches.

Second-pass corrections (F-01 → F-06), applied in this draft:

Finding	Sev.	Resolution in this draft
F-01	Regression High	AI.8a restored (per-hop governance of chained/multi-agent workflows; delegation chain preserved; one ALLOW does not authorize subsequent hops)
F-02	Regression High	AI.12 restored (model/prompt/weights/lineage change = new identity version; triggers EBP scope, capability, threshold, and velocity re-evaluation)
F-03	Regression Medium	8.1 amended to restore the complete delegation chain as an explicit minimum VDR field; term added to Section 2
F-04	Defect Medium	9.2 scoped to E2+ for prevention/detection and given an E1 post-hoc reconciliation floor, closing the mechanism-less MUST at E1
F-05	Disposition accuracy Medium	Annex B G-18 entry corrected to “partially resolved” with the three-element breakdown (vocabulary/neutrality complete; license effective at release; IPR commitment pending)
F-06	Observation Informational	No specification change; certification-sequencing note recorded in Annex E

Annex C — Adversary Model (Normative)

Class	Adversary	Primarily addressed by	In scope at
AC-1	Honest-but-fallible actors (negligence, misconfiguration)	Sections 3–7 (policy, evidence, authority, decisions)	E1+
AC-2	Injected or manipulated AI actors (prompt/input injection, misalignment)	AI profile (AI.4a, AI.6/6a, AI.7, AI.8a, AI.10, AI.12); Section 11	E2+
AC-3	Credentialed insiders acting maliciously with valid authority	6.2, AI.11 (M-of-N independence); Section 8 (VDR); 8.6 (transparency)	E2+
AC-4	Network adversaries (interception, replay, spoofing)	7.6, 7.10, 8.4, 9.3, 13.1	E2+
AC-5	Compromised effectors	9.2, 9.2a, E3 (credential/capability gating)	E3
AC-6	Gate-operator compromise	Section 12, 8.6, 8.8; residual risk MUST be stated in the EBP	E2+ (partial; residual risk explicit)

RF-100.C.1 — Each EBP MUST state which adversary classes are in scope per Governed Action class, consistent with the declared Enforcement Class, and MUST state residual risks for out-of-scope classes. **[I]**

Annex D — Conformance Claims and RF-101 (Normative)

RF-100.D.1 — A conformance claim MUST identify: the EBP version, the Enforcement Class per Governed Action class, the signature scheme in use (8.4), the Gate Availability Class, and the assessment results per Annex D.2. Claims MUST NOT be made against implementations using symmetric MACs as the sole VDR integrity mechanism. **[I]**

RF-100.D.2 (Fault Matrix) — Requirements 7.2 and 7.4 are assessed by fault injection across, at minimum: gate outage; network partition (Actor–gate, gate–Enforcement Point, gate–evidence source); evidence–service failure; authority–service failure; and clock fault beyond the 13.1 drift bound. In each case the assessed behavior MUST be: no execution absent a valid, authenticated, in-window Decision. **[T]**

RF-100.D.3 — The companion document RF-101 (Conformity Assessment Procedures) defines test procedures and the evidence checklist. Until RF-101 v1.0 is published, the skeletal RF-101 checklist published alongside this standard applies. **[I]**

Acknowledgment. Version 1.0 incorporates the full technical gap register and substantial normative drafting input from the structured peer review by James A. Bex (BxTech Systems Solutions LLC, July 2026), and the six second-pass findings (F-01 through F-06) recorded in the July 16, 2026 Second-Pass Review Memorandum, whose contribution the publisher gratefully acknowledges.

Annex E — Second-Pass Findings Disposition (Informative)

This annex records how each finding of the July 16, 2026 Second-Pass Review Memorandum (James A. Bex, BxTech Systems Solutions LLC) is dispositioned in this corrected draft. Proposed normative text supplied by the reviewer was incorporated substantially as offered.

F-01 — Chained-agent per-hop governance (Regression, High). ACCEPTED. Restored as **RF-100:AI.8a**, a distinct requirement alongside the session-attributability requirement (AI.8). Each consequential hop in a chained or multi-agent workflow is independently a Governed Action subject to full verification; the complete delegation chain from the originating Actor is preserved in each VDR; and an ALLOW rendered for one hop does not authorize any subsequent hop.

F-02 — Model/prompt-update re-versioning and EBP re-evaluation (Regression, High). ACCEPTED. Restored as **RF-100:AI.12**, referencing the AI.1 identity fields. A change to an agent’s model, base-model lineage, weights (including fine-tuning), or system prompt constitutes a new agent identity version and MUST trigger re-evaluation of EBP scope, capability set, consequence-threshold assignments, and velocity limits before the updated agent proposes further Governed Actions.

F-03 — Delegation chain in VDR minimum contents (Regression, Medium). ACCEPTED. **RF-100.8.1** amended to restore the complete delegation chain from originating authority to proposing Actor as an explicit minimum VDR field. A “Delegation Chain” definition is added to Section 2 for precision. Together with AI.8a (F-01), this makes the 10.3 reconstruction guarantee mechanically satisfiable from the record schema in delegated and chained scenarios.

F-04 — 9.2 MUST with no mechanism at E1 (Defect, Medium). ACCEPTED (second option: E2+ scoping with the E1 reconciliation floor retained). **RF-100.9.2** rewritten so prevention/detection via 9.2a is scoped to Enforcement Class E2 and above, and Class E1 carries a post-hoc reconciliation floor: the EBP MUST declare a reconciliation procedure comparing effector records against VDR parameter digests on a stated cadence, with divergences recorded and alerted. This removes the mechanism-less, unassessable MUST at E1 while preserving a detection floor there.

F-05 — Annex B overstates the G-18 disposition (Disposition accuracy, Medium). ACCEPTED. The Annex B G-18 entry is corrected from “resolved” to “partially resolved,” stating the three elements explicitly: vocabulary and neutrality complete; CC BY 4.0 effective at release; patent/IPR commitment now published in draft (`IPR-COMMITMENT.md`), pending counsel review, with the binding version to be finalized before any certification program launches.

F-06 — HMAC-only implementations outside conformance (Observation, Informational). NOTED; no specification change. Requirement 8.4 (asymmetric signatures) and D.1 (no conformance claim for sole-symmetric-MAC implementations) are correct and are retained without weakening. The planning consequence is recorded for the publisher: no conformance claim can be made for the

current ExecutionProof platform until asymmetric signing ships, and a certification program (RF-102) cannot credibly launch before the reference implementation itself conforms. Recommended sequencing: (1) 8.4-conformant asymmetric signing in the platform; (2) published patent/IPR commitment (F-05, Section 0.1); (3) RF-101 v1.0; (4) any certification launch.

Freeze status. F-01 through F-05 are dispositioned in normative text in this draft; F-06 requires no specification change. This draft is offered as the release candidate for re-verification against the second-pass memorandum.

© 2026 Remnant Fieldworks Inc. · RF-100 v1.0 DRAFT (unified, second-pass corrected) · Specification text intended for release under CC BY 4.0 (text only; no rights granted to trademarks, service marks, certification marks, software, hosted services, APIs, trade secrets, or patents — see Section 0.1). ExecutionProof™, ProofRecord™, Proof Before Power™, Verification Before Execution™, and Remnant Fieldworks™ are trademarks of Remnant Fieldworks Inc. and do not appear in normative requirement language.